

Convolutional Neural Networks For Image Classification

Sam Gilbert
a1737770

Abstract

Convolutional Neural Networks (CNN) are a type of machine learning model that can be applied to the image classification problem. In this report an investigation was conducted into a subset of the available CNN's and their performance on classifying objects in images from an open source dataset. The impacts of adjusting various hyper parameters on the models overall performance were also investigated and discussed.

1. Introduction

Image classification is the process of taking an image and identifying entities that exist in the image [2]. The applications of image classification are extensive for everything from security cameras to medical imaging. Convolutional Neural Networks (CNN) are a type of Neural Network which is a subset of machine learning [4]. CNN's excel at solving the image classification problem for a variety of reasons including its ability to detect patterns at different levels of complexity. The first proposed implementation was the LeNet-5 architecture proposed in the 1990's. Since then several implementations have been developed and are applied across different problem spaces. With the increased capabilities and availability of Graphics Processing Units (GPU's) and their strength at model training, image processing and classification maturity has increased dramatically.

The CIFAR-10 dataset contains sixty thousand 32x32 colour images containing one of ten object classes [1]. This dataset was used for training and testing the models used through this report. The dataset contains 50,000 images for training and 10,000 images for testing. Each image in this dataset contains a single type of entity from a set of types. The ability of various models to categorise the type of object in was investigated and summary statistics of each model was compared to determine the strengths and weaknesses of the models. The set of models investigated are;

- VGG-16
- VGG-19

- ResNet-18
- ResNet-34
- ResNet-50

the differences between the models will be investigated in the Models section of this report.

2. CNN

CNN's have three main types of layers, convolutional, pooling and fully connected.

2.1. Convolutional Layer

The convolutional layer is responsible for feature extraction and work by applying a defined filter over a set MxM pixels from the input image. This MxM window is "slid" across the image and the dot product for each set of pixels is calculated with the dot product being defined as;

$$a \cdot b = \sum_{i=1}^m a_i b_i \quad (1)$$

The output from this convolutional layer is known as a feature map. This feature map is a 2 or 3 dimensional matrix of values that reflects the presence of certain features in a given image such as edges or patterns. A convolutional layer often applies multiple filters and returns multiple feature maps. The number of feature maps returned is equal to the number of filters that were applied in the given layer. At the early stages of the CNN the convolutional layers and the specified filters mainly focus on lower level components of the input image as mentioned above. The convolutional layers later in the network are more abstract and are used to capture higher level features of the image such as eyes or wheels. Often the feature maps generated are passed through a Rectified Linear Unit (ReLU) which is defined as;

$$f(x) = \max(0, x) \quad (2)$$

This ReLU introduces non linearity into the model allowing complex patterns and repetitions. It also introduces sparsity into the model through the zeroing out of negative values in the maps increasing the overall efficiency of the model.

2.2. Pooling Layer

The pooling layer in a CNN is responsible for reducing the dimensionality of the convolutional layers feature map output. It performs a downsampling on a given feature map similar to the convolutional layer by applying a function over a sliding window section of each feature map. A function is then applied to each window and the combined output is passed on to the next layer. This process increases efficiency in the model, increases adaptability and reduces the number of parameters the model needs to learn reducing the possibility of overfitting. By reducing the dimensionality of the feature maps, the model becomes less sensitive to specific positions or orientations of given features and becomes more effective at detection. The two most common pooling methods are max pooling and average pooling. Given an $N \times N$ window from a feature map, max pooling returns the maximum value present in that window whereas average pooling returns the average value of the given window.

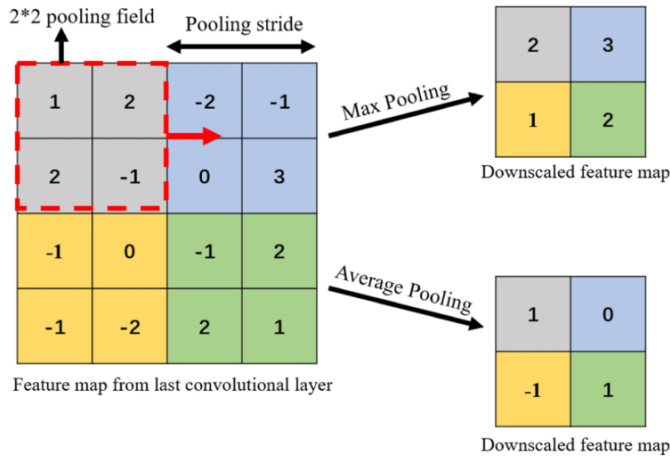


Figure 1. Example of max and average pooling on output of previous convolutional layer [6]

2.3. Fully Connected Layer

The final layer or layers in a CNN is the fully connected layer. These fully connected layers have each input connected to each output by a learnable rate. The final fully connected layer typically has the same number of output nodes as the number of classes [10]. This final output layer then includes an activation function which is typically different than the output function in the convolutional layers. This output function is set dependant on the task the CNN model is trying to achieve. In the case of the image classification problem a softmax activation function is typically used [8] since it can be used to determine what the image is

most likely to contain. The softmax function is defined as follows;

$$\sigma(\vec{z})_i = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}} \quad (3)$$

where σ is the softmax, \vec{z} is the input vector, e^{z_i} is the standard exponential function for input vector, K is the number of classes and e^{z_j} is the standard exponential function for output vector. The output vector generates a set of probabilities for each class in K with the max being selected as the final classification of the image.

2.4. Training

Training the CNN is done through a method known as backpropagation. Backpropagation dates back to the 1960's and was popularised in a paper published in 1986 [9]. Backpropagation is an iterative algorithm that aims to minimise the cost function through two passes of the CNN model. The forward pass operates as it would during the testing phase with raw inputs being provided to the input layer of the CNN. Then the data traverses the network and the final output is returned. This output is compared to the actual values in the training data and an error value is determined. A common method of calculating the error is the mean squared error (MSE) which is defined as;

$$MSE = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2 \quad (4)$$

where n is the number of observations, Y is the observed values and \hat{Y} is the predicted values.

Once this MSE has been calculated the backwards pass takes this calculated error and propagates it back through the network. The backwards pass utilises a gradient descent method to determine the gradient for each weight and bias in the network layers and is used to inform the magnitude of change that can be applied to minimise the error in the next forward pass. By utilising the chain rule, calculating this change of parameters with the goal of minimising the error can be done efficiently. This two pass process is done a predefined number of times, often called the epoch, with the final weights and biases for each layer being the values used in the model for the testing phase. Determining an optimal value for this epoch will form the basis of investigation through this report since a balance must be struck between accuracy improvements and training time required. Another concern for the CNN model is the possibility of overfitting and underfitting. This is where the model contains too many layers and learns the optimal values for the training set. This results in high accuracy when training however could result in underperformance on new data during testing. If this occurs it can be a sign that there are too

many convolutional layers in the CNN and improved performance may be observed with a simpler network. Since this report uses pre designed networks the risk of this is minimal however it needs to be monitored when investigating the impact of parameter changes.

3. Pytorch

Although it is possible to build a CNN from scratch, open source tools have allowed for rapid development and implementation of these models. Pytorch is one of these tools which is available in the Python programming language. Pytorch enables rapid prototyping of different models through its publishing of base implementations on a large number of common machine learning models. For the models investigated in this report, a base Pytorch implementation was available for each of them. These models will be described in section 4 of this report. Pytorch exposes low level api calls that allow the user to modify the models which greatly aids in the types of investigation and experimentation that are possible. Pytorch can be used for model training and testing and provides robust functionality for the end to end process of defining and then using a given machine learning model. Pytorch has native support for running these models on GPU's through its Cuda integration. Cuda is a parallel computing platform and programming model developed by NVIDIA [7] that allows developers to copy data to supported graphics cards and utilise the highly parallel architecture inherent to GPU's to increase training time greatly. Since the training process involves large numbers of operations to be done on matrices, this work can be parallelised effectively for massive performance gains.

4. Data

As mentioned in the introduction, the dataset used for this report was the CIFAR-10 image dataset [1]. The dataset is comprised of 60,000 32x32 colour images that contain a single entity of one of ten classes. These classes are airplane, automobile, bird, cat, deer, dog, frog, horse, ship and truck. This dataset is split into a training and a testing set with 50,000 of the images being used for training and the remaining 10,000 being used for testing

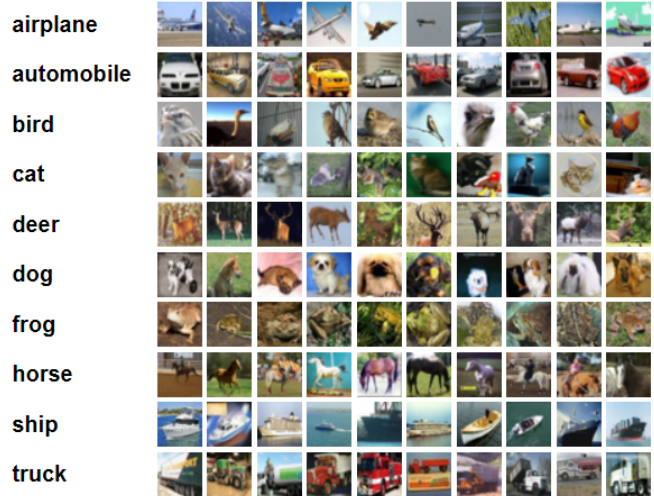


Figure 2. Examples of each class in the CIFAR-10 image dataset

5. Models

Due to the maturity of CNN's and their application to the image classification problem, a large number of models have been developed specifically for this application. Implementing a subset of these models will form the basis of the investigation on how modifying training parameters and training epochs can influence overall performance. The base implementation available in Pytorch will be used.

5.1. VGG-16

VGG-16 contains three fully connected layers and 13 convolutional layers [5]. It also contains five pooling layers that are placed consistently throughout the network. There is a pooling layer after the first two convolutional layers and the third and fourth convolutional layer. Then there is a pooling layer after each set of three convolutional layers for the rest of the network. This is then followed by the three fully connected layers which provide the final output for the model. This architecture results in 16 layers with learnable parameters.

5.2. VGG-19

VGG-19 is extremely similar to VGG-16 however an additional convolutional layer is added in the last three blocks. This results in an architecture of two convolutional layers, one pooling layer repeated twice. Then four convolutional layers, one pooling layer repeated three times. Then finally the three fully connected layers at the end of the network. VGG-19 has three additional layers with learnable parameters.

5.3. ResNet-18

The ResNet architecture differs from traditional neural networks. The ResNet-18 model contains a single convolutional layer initially then two pooling layers and eighteen convolutional layers between the pooling layers before a single fully connected layer. Traditionally when the network depth increases, accuracy gets saturated [3]. The ResNet architecture aims to address this performance degradation by introducing a deep residual learning framework. This is achieved by implementing a residual block which contains a set of convolutional layers and a skip connection. Instead of learning a mapping, $H(x)$, to be fit by the stacked layers, the block learns $F(x) = H(x) - x$ or $H(x) = F(x) + x$ where x is the input to the first layer in the set. The output of the residual block is;

$$Output = F(x, weights) + x \quad (5)$$

where $F(x, weights)$ is the learned residual. The skip connections perform identity mapping, and their outputs are added to the output of the stacked layer [3]. This allows the backpropagation process to effectively "skip" layers by focussing on the learned residual.

5.4. ResNet-34

The ResNet-34 model is functionally the same as the ResNet-18 model however instead of having eighteen convolutional layers between the pooling layers there are thirty four.

5.5. ResNet-50

As the ResNet architecture increases the number of convolutional layers a more efficient "bottleneck" design is used for the residual block. Each block is made of a 1x1, 3x3 and 1x1 convolution [3]. The 1x1 blocks are made for dimensionality reduction and expansion to the original dimensions. The 3x3 block is for feature extraction as normal. The skip connections are added between the start and end of each of these blocks. The ResNet-50 model replaces each two layer block present in the ResNet-34 network with the bottleneck design resulting in a fifty layer network.

6. Method

CNN's have a number of parameters that can be set to alter the models final output and corresponding performance. An investigation into a subset of these was conducted and the results discussed. The key parameters shared across the investigated models were backpropagation epoch, learning rate and batch size. For all the models a baseline implementation was built and its performance quantified. The metrics used to evaluate performance was accuracy, precision, recall and the F1 score. These were calculated through the confusion matrix generated after running the test dataset

through the model. In a multi class classification problem such as image classification, the generation of a confusion matrix becomes more complex than the binary classification space. The confusion matrix is an $n \times n$ matrix with n being the number of classes. Each entity type has a row and column and the metrics are calculated on a per entity type basis. To retrieve the overall values for the metrics the mean of each is calculated and returned as the final result.

Accuracy calculates what proportion of the classifications were correct and is defined as;

$$Accuracy = \frac{P_{correct}}{P_{correct} + P_{incorrect}} \quad (6)$$

where P is the prediction made.

Precision is a measure of the correctness of positive predictions which in the multi class image classification problem can answer the question out of all the predictions made, what proportion were actually the predicted class? This can be defined as;

$$Precision = \frac{1}{n} \sum_{i=0}^n \frac{M_{ii}}{\sum_j M_{ji}} \quad (7)$$

Where M is the confusion matrix, n is the number of classes

Recall is a measure of the proportion of correct predictions for a class out of all the cases where that specific class. It is a measure of the number of times the model correctly identifies a given class from all the actual instances of that class in a dataset. This can be defined as;

$$Recall = \frac{1}{n} \sum_{i=0}^n \frac{M_{ii}}{\sum_j M_{ij}} \quad (8)$$

The F1 score is the harmonic mean of precision and recall and is defined as;

$$F1Score = \frac{2 * Precision * Recall}{Precision + Recall} \quad (9)$$

Given the time it takes to train a new model on each parameter change, the impacts to performance will be performed on one of each type of model. For this investigation the VGG-16 and ResNet-18 models were investigated since they are the models of each type with the smallest number of convolutional layers and should provide representative performance changes as a result of the hyper parameter investigations.

To investigate the effect different epochs have on model performance during the backpropagation process a set of values will be investigated. The baseline epoch used for all the models is 10 and the values 1,3,5,7,10,15,20,40 and 80 will be tested. Similarly the learning rate will be tested with a set of value with a baseline of 0.001. The values tested will be, 0.1, 0.01, 0.001 and 0.0001. Finally the batch size

will be investigated with a baseline of 32 and the values investigated will be 8,16,32 and 64.

As the pretrained models are being used and they have been trained more effectively than the hardware available to this project, freezing of the model parameters will be done. This involves not updating the parameters except in the final fully connected layer which needs to be done to only predict one of the ten classes in the CIFAR-10 dataset. This will greatly increase the speed of training and will allow the leveraging of the well trained models capabilities. It can also reduce the risk of overfitting as if the model was solely trained on the CIFAR data then it may learn features inherent to it reducing the models generalisation abilities. Loss will be calculated through cross entropy loss, also known as log loss and for the multiclass classification can be defined as,

$$Loss(\hat{y}, y) = - \sum_k^K y^k \log \hat{y}^k \quad (10)$$

where y^k is 0 or 1 indicating whether the class label k is the correct classification [11].

Finally the images will be preprocessed through resizing them to a 224x224 image which are the dimensions that the underlying models being used were trained on. Then normalizing them using the mean and standard deviations of the images that the pretrained models were trained on. This will allow the model to be further trained on images that are similar in shape and scale to which they were initially trained on thus theoretically improving model performance

7. Results

7.1. Baseline

As discussed in the method section, the first step of the investigation into the effects of modifying the hyper parameters a set of baseline results must be obtained. With epoch = 10, learning rate = 0.001 and batch size = 32 Table 7.1 shows the gathered metrics.

Model	Accuracy	Precision	Recall	F1 Score
VGG16	0.870	0.875	0.870	0.872
VGG19	0.882	0.885	0.882	0.883
ResNet18	0.903	0.904	0.903	0.902
ResNet34	0.907	0.910	0.908	0.908
ResNet50	0.892	0.894	0.892	0.892

Table 1. Baseline metrics for all investigated models

It can be observed that the best performing model out of the VGG models is VGG19 with the highest values for all computed metrics. For the ResNet models, ResNet34 performed the best again with the highest metrics across the board.

7.2. Epoch Investigation

Determining an optimal or near optimal epoch value is important as it influences future hyper parameter investigations. Due to the training time required it is infeasible to test every combination of each hyper parameter as the number of models to be trained increases exponentially as the search space expands. As such each hyper parameter will be investigated in isolation. The impacts of this on the validity of the results will be discussed in the Discussion and Future Work sections of the report. The results of the epoch investigation are outlined in Table 2

Epoch	Accuracy	Precision	Recall	F1 Score
1	0.842	0.852	0.842	0.843
3	0.861	0.868	0.861	0.862
5	0.864	0.872	0.864	0.866
7	0.871	0.877	0.871	0.872
10	0.870	0.875	0.870	0.872
15	0.862	0.883	0.862	0.866
20	0.874	0.880	0.874	0.875
40	0.868	0.882	0.868	0.873
80	0.860	0.884	0.861	0.867

Table 2. Epoch results for VGG16

Interestingly the accuracy values are almost identical to the recall results. This occurs when the dataset is balanced, has a roughly even number of each class. This will be investigated further in the Discussion section. These results are plotted in the below image;

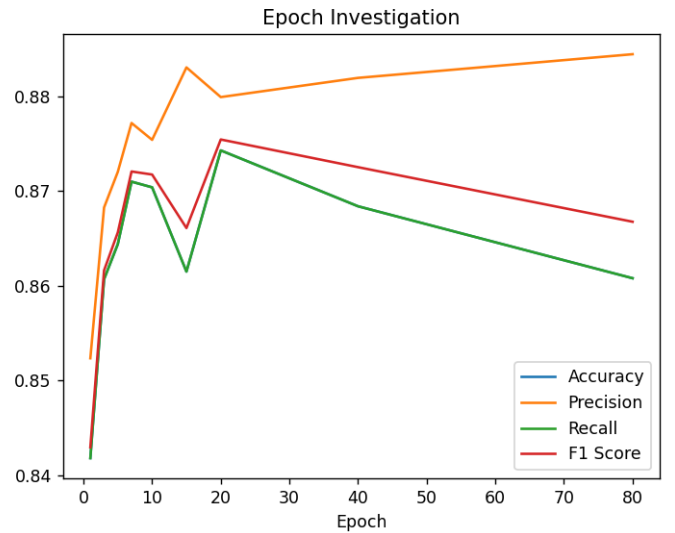


Figure 3. Plotted output from Table 2

Plotting the same set of epochs for the ResNet18 model yields the plot shown in Figure 4.

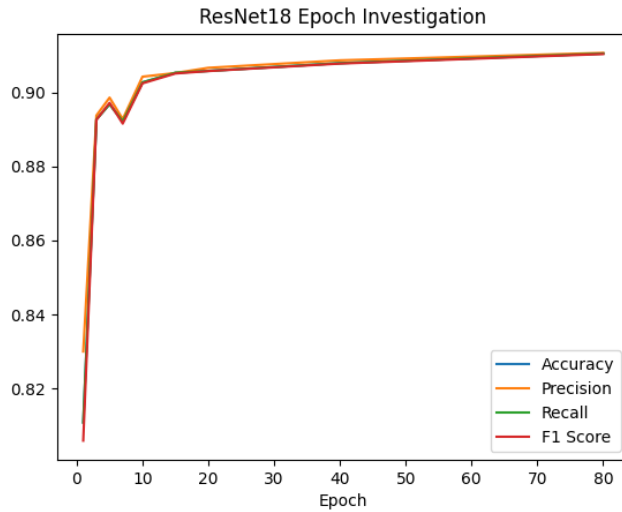


Figure 4. ResNet18 epoch investigation results

Unlike the VGG16 model, the metrics continue to increase as the number of epochs investigated increases. The values investigated clearly indicate the optimal value has not been found however the amount of performance increase doesn't appear to change much for epochs greater than 20. Balancing the time taken to train each new model with epochs greater than the upper bound used, 80, and the corresponding increase to performance, an epoch value of 20 has been found to be a suitable value.

7.3. Learning Rate Investigation

When investigating the impact learning rate has on the overall model performance, the results from the epoch investigation were used as a baseline. That is the number of epochs used was 20 as that was found to have the best performance. As discussed in the Method section, the investigated values were 0.1, 0.01, 0.001, 0.0001, 0.00001 and 0.000001. The metrics were found and are displayed in Table 3 and were plotted in figure 5.

LR	Accuracy	Precision	Recall	F1 Score
0.1	0.100	0.110	0.100	0.018
0.01	0.100	0.110	0.100	0.019
0.001	0.874	0.880	0.874	0.875
0.0001	0.887	0.889	0.887	0.887
0.00001	0.895	0.896	0.895	0.895
0.000001	0.878	0.878	0.878	0.877

Table 3. Learning rate results for VGG16

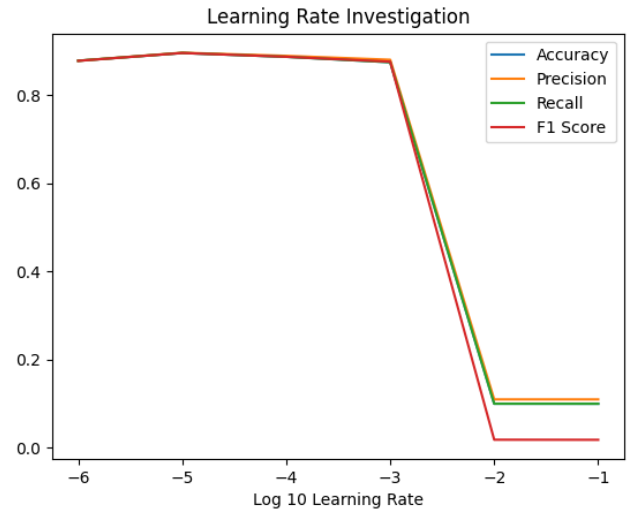


Figure 5. Plotted output from Table 3

7.4. Batch Size Investigation

As mentioned in the method section of the report the investigation into the batch size's impact of model performance was done using batch sizes of 8, 16, 32 and 64. The theoretical impact of altering the batch size will be discussed in the Discussion section of this report. For this investigation the VGG16 model was once again used with a learning rate of 0.00001 and 20 epochs. Setting this learning rate and epoch statically may not be an optimal solution however this will also be discussed later.




Figure 6. Batch size investigation results

From figure 6 it can be observed that as the batch size increases, the models performance trends upwards.

8. Discussion

CNN's expose a number of parameters that can be modified to impact overall performance. Throughout this report a subset of these have been chosen and investigated to measure each of their impacts. The epoch value determines the number of iterations of backpropagation that occurs during the training process. Increasing the number of epochs allows the model more iterations for optimising the gradient to the training dataset. More opportunities for optimising feature extraction and detection within the model are provided however a balance needs to be struck to avoid overfitting. Overfitting is where the model becomes too optimised at maximising performance on the training data and reduces the performance when new data is presented to the model. The epoch investigation on the VGG16 model showed that model performance when measured by recall, f1 score and accuracy peaked at an epoch value of 20. The precision metric continued increasing as the epoch value increases which implies that the models performance on its positive predictions continues to increase. That is the number of true positives compared with the total number of positive predictions continues to increase. This is paired with a decrease in the models recall which incorporates false negatives in its calculation and implies that the model is generating more false negatives as the epoch value increases. A different observation was made when the ResNet18 model was investigated as the metrics continued to increase as the epoch value exceeded the upper bound tested, 80. This implies that the optimal value for the epoch was not found in this report and in the future work section this will be expanded on. The impact of selecting the right epoch is drastic with the performance difference between an epoch of 1 and the optimal value for the VGG16 model being approximately 3.2% and with the ResNet18 model the difference being almost 10%.

The learning rate controls the rate at which the parameters at each layer are updated throughout the training process. A higher learning rate results in a greater rate of change. If the learning rate is set too high then the model will likely reach its optimal performance quicker however the optimal parameter values found may overshoot the models actual optimal values resulting in degraded performance. This is caused by the model "overshooting" the optimal values and the learning rate used not providing the granularity required. This needs to be balanced against the number of epochs required for the corresponding learning rate to converge on the optimal parameter values. With a smaller learning rate the number of iterations will be increased. From the investigation conducted into how

the learning rate impacts the models overall performance, if the clearly unsuitable values are ignored, 0.1 and 0.01, the optimal learning rate can yield another approximately 2% performance increase.

The batch size investigation suggested that as the batch size increased, the model performance also increases. By increasing the batch size the backpropagation process utilises more of the training data resulting in smoother gradient updating. It does come with the risk of overfitting the data as this larger rate of change of the parameters runs the risk of getting stuck at local minima or maxima values resulting in not returning the optimal values. The size of the batches used also relies on the learning rate being appropriately tuned. For the investigation conducted in this report the learning rate was statically set. A smaller batch size relies on a smaller learning rate otherwise there is a risk of missing the optimal solution. On the other hand, a larger batch size can tolerate a larger learning rate however it still needs to be effectively tuned to avoid poor convergence and generalisation.

A potential issue that could have impacted the results is that for this investigation, the base models used before training was the pretrained version published by Pytorch. Since the cifar10 dataset is a commonly used dataset, it is safe to assume that the dataset was used in the pretrained model. Performance was found to be significantly better after training the models on the cifar10 dataset. Each of the models investigated used the pretrained versions therefore providing a consistent base for the investigations conducted. Future work to repeat these investigations without the pretrained base will be expanded in the future work section.

9. Future Work

For future work on this project, further investigations will be done on finding the optimal epoch value for the ResNet18 models. As mentioned in the discussion, for the epochs that were tested the models performance was still increasing. This implies that the optimal epoch has not been found. Continuing to train ResNet18 models with increasing epochs to identify the inflection point will be done. Throughout this report the VGG16 and ResNet18 models were used for the various hyper parameter investigations, assumptions have been made that the findings for these models will be consistent with the different forms of the models. An investigation into this assumption will be completed by training each form of the models on the found optimal values for the parameters. Then determining that the performance gains found are replicated in the other models.

Only a subset of available CNN's were investigated in this report. Expanding this subset will be investigated in fu-

ture work. Some of the other models to investigate are the AlexNet and GoogLeNet architectures. Re-running the investigations on the base versions of the models made available in Pytorch will be conducted. Determining the impact that the pretrained model has on model performance for the cifar10 dataset will be identified. Determining if model overfitting has tainted the results will need to be found and model performance using the found optimal values will be investigated. It is expected that performance will be reduced as the pretrained versions of the models have undergone a robust training process on a huge dataset. This report has ignored the time portion of model training, however this is a key consideration when determining optimal model parameters. Given an infinite time horizon the absolute optimal values can be determined. Future work will include investigating the impact that modifying parameters has on the overall time taken for training the model.

10. Appendix

References

- [1] Will Cukierski. Cifar-10 - object recognition in images, 2013.
- [2] Weili Fang, Peter E. D. Love, Hanbin Luo, and Lieyun Ding. Computer vision for behaviour-based safety in construction: A review and future directions. *Advanced Engineering Informatics*, 43:100980, 2020.
- [3] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition, 2015.
- [4] IBM. What are Convolutional Neural Networks? — IBM — ibm.com. <https://www.ibm.com/topics/convolutional-neural-networks>. [Accessed 12-10-2024].
- [5] Ali Hasan Md. Linkon, Md. Mahir Labib, Tarik Hasan, Mozammel Hossain, and Marium-E-Jannat. Deep learning in prostate cancer diagnosis and gleason grading in histopathology images: An extensive study. *Informatics in Medicine Unlocked*, 24:100582, 2021.
- [6] Liang Minfei, Yidong Gan, Ze Chang, Zhi Wan, Erik Schlangen, and Branko Šavija. Microstructure-informed deep convolutional neural network for predicting short-term creep modulus of cement paste. *Cement and Concrete Research*, 152:106681, 02 2022.
- [7] Fred Oh. What is cuda?, Jan 2022.
- [8] Jingli Ren and Haiyan Wang. Chapter 3 - calculus and optimization. In Jingli Ren and Haiyan Wang, editors, *Mathematical Methods in Data Science*, pages 51–89. Elsevier, 2023.
- [9] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. Learning representations by back-propagating errors. *Nature*, 323(6088):533–536, Oct. 1986.
- [10] Rikiya Yamashita, Mizuho Nishio, Richard Kinh Gian Do, and Kaori Togashi. Convolutional neural networks: an overview and application in radiology. *Insights into Imaging*, 9(4):611–629, June 2018.
- [11] Iffat Zafar, Giounona Tzanidou, Richard Burton, Nimesh Patel, and Leonardo Araujo. Hands-on convolutional neural networks with tensorflow.